

PROCESS SYNCHRONIZATION II

Ruben Acuña

Spring 2019



SEMAPHORES (5.6)

SEMAPHORES

- Generalization over mutex to support more than one resource.
 - A more mutex is about excluding an action (running code in a second process) while a semaphore is about acquiring some unique resource from a set.
 - Semaphores end up working a little like malloc/free.
- Let some integer S be called a semaphore value, and define two operations:
 - Wait(): blocks until S is positive
 - Signal(): increases S by 1.
 - Both should be atomic operations.

```
wait(S) {  
    while (S <= 0);  
    S--;  
}  
  
signal(S) {  
    S++;  
}
```

USAGE

- Two samples:
 - Solving the critical-section problem.
 - Enforcing execution order.

//critical section problem

//shared data

semaphore lock = 1;

```
do {  
    wait(&lock);  
    // critical section  
    sign(&lock);  
    // remainder section  
} while (true);
```

What would happen
if this was 0?

//enforcing execution order

//shared data

semaphore lock = 0;

```
void proc1() {  
    printf("line 1");  
    signal(&lock);  
}
```

```
void proc2() {  
    wait(&lock);  
    printf("line 2");  
}
```

IMPLEMENTATION

- Say we want to implement a semaphore S where $S \in I$, and want to avoid busy waiting.
- The solution is relatively simple, use a list to store process and block them instead of looping.

```
typedef struct {
    int value;
    struct process *list;
} semaphore;

wait(semaphore* S) {
    S->value--;
    if (S->value < 0) {
        //add this process to S->list
        block();
    }
}

signal(semaphore* S) {
    S->value++;
    if (S->value <= 0) {
        //remove a process P from S->list
        wakeup(P);
    }
}
```

TERMINOLOGY

- Consider the case of processes P_1 , and P_2 . Both of these processes need access to resources controlled by semaphores S_1 , and S_2 . If the processes lock a (different) semaphore in parallel, then both will be unable to lock the second semaphore and go into *deadlock*.
- We also use the term *starvation* to refer to a process that gets "stuck" on a semaphore and cannot proceed.
- Consider the case of a some process P_1 , that requires a resource used by P_2 . Let P_1 be of a much higher priority than P_2 . Then, P_1 may wait because P_2 isn't given a chance to execute and release the resource. This is called *Priority Inversion*. Inversion in the sense that the lower priority process should now be considered higher priority.
- One solution is *priority-inheritance*, which basically says if a resource is needed by a high priority process, then other users of that resource should inherit the high priority.



CLASSIC PROBLEMS OF SYNCHRONIZATION (5.7)

THE BOUNDED BUFFER PROBLEM

- Let's try rephrasing the producer-consumer problem (over a bounded buffer), using semaphores.
- Going to have four variables:

```
//shared data
int n;
//also a buffer data structure
semaphore buf_mutex = 1;
semaphore empty = n;
semaphore full = 0
```

```
//producer
do {
    // produce next_produced
    wait(empty);
    wait(buf_mutex);
    // add next_produced to the buffer
    signal(buf_mutex);
    signal(full);
} while (true);

//consumer
do{
    wait(full);
    wait(buf_mutex);
    // move from buffer to next_consumed
    signal(buf_mutex);
    signal(empty);
    // consume next_consumed
} while (true);
```


THE READER-WRITERS PROBLEM

- New problem: considering a single source of data that multiple threads may read or write into. Multiple reads are fine, but when a write is happening, the data should be constant.
- There are a couple of design approaches – here we'll prioritize the readers of the writers. (Means that writers may starve).
- Need three variables:

```
//shared data
semaphore rw_mutex = 1;
semaphore rc_mutex = 1;
int read_count = 0;
```


```
//writer process
do {
    wait(rw_mutex);
    /* writing is performed */
    signal(rw_mutex);
} while (true);

//reader process
do {
    wait(rc_mutex);
    read_count++;
    if (read_count == 1)
        wait(rw_mutex);
    signal(rc_mutex);
    /* reading is performed */
    wait(rc_mutex);
    read_count--;
    if (read_count == 0)
        signal(rw_mutex);
    signal(rc_mutex);
} while (true);
```

THE DINING PHILOSOPHERS

- Consider the following scenario: you have 5 philosophers sitting around a table, on the table are five plates of food, and five chopsticks.
- Occasionally, a philosopher gets hungry, in which case they must acquire two chopsticks, eat, and then puts those chopsticks down.
- Clearly there aren't enough chopsticks for everyone to eat at once...
- How well would the following solution fair in feeding philosophers?

```
//Philosopher i
do {
    wait(chopStick[i]);
    wait(chopStick[(i + 1) % 5]);
    // eat
    signal(chopStick[i]);
    signal(chopStick[(i + 1) % 5]);
    // think
} while (1);
```





MONITORS (5.8)

USING SEMAPHORES

- Consider the three mini-programs to the right. For each case, picture what will happen if we run threads on this piece of code.
- What will happen with program A? B? C?
- Next, we'll look at monitors which is an OOP-style approach to blackboxing the mutual exclusion of threads that we are after.

```
//program A  
signal(mutex);  
//critical section  
signal(mutex);
```

```
//program B  
signal(mutex);  
//critical section  
wait(mutex);
```

```
//program C  
wait(mutex);  
critical section  
wait(mutex);
```

MONITOR CONCEPT

- Assume that we have some class that is a so-called “monitor”.
- Then, the methods in that class all have access to the instance variables, but only one process at a time is allowed to run within the monitor class.

```
monitor class Account {  
    int balance;  
  
    Account(int opening) {  
        balance = opening;  
    }  
  
    void deposit(int amount) {  
        balance = balance + amount;  
    }  
  
    void withdraw(int amount) {  
        balance = balance - amount;  
    }  
};
```

CONDITIONAL VARIABLES

- In addition, we'll introduce a new idea to alter program flow: conditional variables.
- A conditional variable exists within a monitor, and is acted upon by two functions (wait, and signal) so that processes using that monitor continue in certain way.
- When a process calls wait, it will block until the conditional is ready.
- When a process calls signal, then (typically) the monitor will transfer control to one of the processes currently waiting on that condition.
 - See signal-and-wait vs signal-and-continue.


```
monitor class BoundedBuffer {  
    //shared data  
    int buffer[MAX];  
    int fill, use;  
    int fullEntries = 0;  
    cond_t empty;  
    cond_t full;  
  
    void produce(int element) {  
        if (fullEntries == MAX)  
            wait(&empty);  
        buffer[fill] = element;  
        fill = (fill + 1) % MAX;  
        fullEntries++;  
        signal(&full);  
    }  
  
    int consume() {  
        if (fullEntries == 0)  
            wait(&full);  
        int tmp = buffer[use];  
        use = (use + 1) % MAX;  
        fullEntries--;  
        signal(&empty);  
        return tmp;  
    }  
}
```

BACK TO DINING

```
monitor DiningPhilosophers {  
    enum { THINKING; HUNGRY, EATING) state[5];  
    cond_t self[5];
```

```
void pickup(int i) {  
    state[i] = HUNGRY;  
    test(i);  
    if (state[i] != EATING)  
        self[i].wait();  
}  
  
void putdown(int i) {  
    state[i] = THINKING;  
    // test left and right neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```

```
}
```



```
void test(int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING)) {  
        state[i] = EATING;  
        self[i].signal();  
    }  
}  
  
void initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

IMPLEMENTING A MONITOR USING SEMAPHORES

- Each monitor includes a class wide mutex semaphore that processes must use to access functions.
- As an alternative to a mutex, processes actively in a monitor may “pause” by signaling mutex and instead waiting on next.
 - next_count is number of waiting processes.
- Each conditional variable x has a semaphore (_sem) and int (_count).

```
F() {  
    wait(mutex);  
  
    //body of F  
  
    //continue next process  
    if (next_count > 0)  
        signal(next);  
    else  
        signal(mutex);  
}
```

```
//x.wait()  
x_count++;  
if(next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x_count--;
```

```
//x.signal()  
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```




SYNCHRONIZATION EXAMPLES (5.9)

Skipping 5.9.1 (Synchronization in Windows), and 5.9.3
(Synchronization in Solaris)

LINUX

- There are various mechanisms in the Linux kernel for providing low-level synchronization:
 - Atomic integer operations (arch/*/atomic.h):
 - `typedef atomic_t`
 - `int atomic_set(atomic_t *a, int value);`
 - `int atomic_add(int value, atomic_t* a)`
 - `int atomic_sub(int value, atomic_t* a)`
 - `int atomic_inc(&counter)`
 - `int atomic_read(atomic_t *a)`
 - Spinlocks (`spin_lock_irqsave`, `spin_lock_irqsave`)
 - Mutex support (`kernel/locking/mutex.c`):
 - `int mutex_init(mutex_t *mp, int type, void * arg);`
 - `int mutex_lock(mutex_t *mp);`
 - `int mutex_unlock(mutex_t *mp);`

PTHREADS

- Syntax very straight forward.

```
#include <pthread.h>
pthread_mutex_t mutex;

// create the mutex lock
pthread_mutex_init(&mutex, NULL);

// acquire the mutex lock
pthread_mutex_lock(&mutex);

// critical section

// release the mutex lock
pthread_mutex_unlock(&mutex);
```

```
#include <semaphore.h>
sem_t sem;

// Create the semaphore and set it to 1
sem_init(&sem, 0, 1);

// acquire the semaphore
sem_wait(&sem);

// critical section

// release the semaphore
sem_post(&sem);
```



ALTERNATIVE APPROACHES (5.10)

SOLVING THE CRITICAL SECTION

- Transactional Memory: the basic idea of transactions emerged from databases (finance) where need to ensure that all of set of operations happens at once (i.e., atomically).
- Thinking back to the discussion on automatic threading, OpenMP also supports marking something as a critical region.
- Most of our problems come from mutable state – if we do a way with state, then many problems vanish... so, just use *functional style programming*! All data will be immutable.

```
#pragma omp parallel
{
    //stuff in parallel

    #pragma omp critical
    {
        printf("critical");
    }
}
```